




103 GB database, 50-read workload. DRAM share drops from 46% (no cache) to 24% (4-way) - the cache intercepts exactly the expensive misses. X-axis width = the SOFTWARE S2 lookup-cache's associativity. The L2/LLC/DRAM breakdown itself is REAL Luna HARDWARE cache behavior (L1d 48KB/12-way, L2 2MB/16-way, L3 105MB/15-way) - this chart shows how a software design choice changes what the real hardware caches actually see.