


50 reads 2000 reads

**Same 50MB DB, 50 reads vs 2000 reads:
the cache-hurts result already existed, and holds**



Both bars use the SAME 50MB DB (sample_targeted). The cache was already worse than no-cache at just 50 reads (fair batch, 2026-09-09) - this was NOT newly discovered by the 2000-read job. What the 2000-read job adds: ruling out "needs more reads to warm up" as an excuse - the negative result holds (and worsens slightly) at 40x more reads. This says nothing about the 8GB/16GB/103GB DBs, where cache DOES help (~2x) at 50 reads - that comparison has not been run at large read counts and remains open.